

# Containers

for

## Scientific Computing and Data Science

*Tommaso Buvoli*

*September 29, 2020*

# Common Software Problems



# Common Software Problems



**Each project may require many software dependencies.**

# Common Software Problems



**Each project may require many software dependencies.**

- No reproducibility across machines or across time.

# Common Software Problems



**Each project may require many software dependencies.**

- No reproducibility across machines or across time.
- Packages end up littering your hard drive for years even after you've completed a project.

# Common Software Problems



**Each project may require many software dependencies.**

- No reproducibility across machines or across time.
- Packages end up littering your hard drive for years even after you've completed a project.
- Certain software can also be difficult to install.

# Why Containers?

# Why Containers?

1. **Containers are reproducible.**



# Why Containers?

## 1. Containers are reproducible.

- You can save a compute environment and replicate it on any machine at a later date.

# Why Containers?

## 1. Containers are reproducible.

- You can save a compute environment and replicate it on any machine at a later date.
- You can share containers across teams.

# Why Containers?

## **1. Containers are reproducible.**

- You can save a compute environment and replicate it on any machine at a later date.
- You can share containers across teams.

## **2. Containers bundle your software dependencies.**

# Why Containers?

## 1. Containers are reproducible.

- You can save a compute environment and replicate it on any machine at a later date.
- You can share containers across teams.

## 2. Containers bundle your software dependencies.

- No clutter on your local machine.

# Why Containers?

## 1. Containers are reproducible.

- You can save a compute environment and replicate it on any machine at a later date.
- You can share containers across teams.

## 2. Containers bundle your software dependencies.

- No clutter on your local machine.
- Save your containers online and reuse them at a later date.

# A Computer in a Computer



# A Computer in a Computer



- Containers let you run a “*computer in a computer.*”

# A Computer in a Computer



- Containers let you run a “*computer in a computer.*”
- Similar in spirit to virtual machines.



# Container Basics

# Container Basics

1. A **container image** describes the Linux operating system and the software installed in it.

# Container Basics

1. A **container image** describes the Linux operating system and the software installed in it.
  - You can think of an image as a description for a virtual computer.

# Container Basics



1. A **container image** describes the Linux operating system and the software installed in it.
  - You can think of an image as a description for a virtual computer.
2. A **container** is a collection of processes that run from an image.

# Container Basics

1. A **container image** describes the Linux operating system and the software installed in it.
  - You can think of an image as a description for a virtual computer.
2. A **container** is a collection of processes that run from an image.
  - Multiple containers can be run from the same image.


# Containers

Image



# Containers

## Image




**Remark:** images are labeled as name:tag (e.g. myimage:v10)




# Containers

## Image



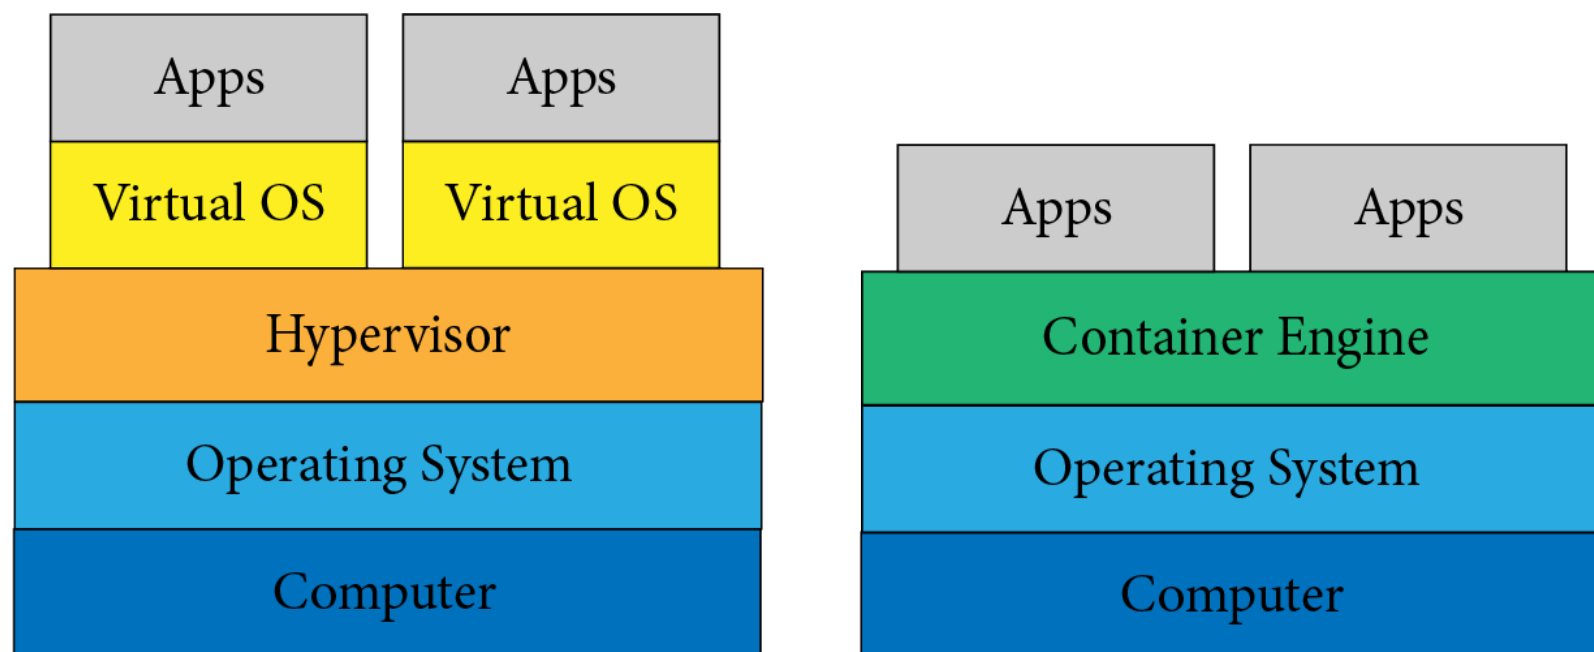
**Remark:** images are labeled as name:tag (e.g. myimage:v10)



**Remark:** Containers are ephemeral. The image is never changed.




# Containers vs Virtual Machines




# Containers vs Virtual Machines

1. Containers are faster to start than virtual machines; no hypervisor, the kernel is shared with host the OS.



# Containers vs Virtual Machines

1. Containers are faster to start than virtual machines; no hypervisor, the kernel is shared with host the OS.
2. It is much simpler to share and distribute container images than VM images.



# Three Container Engines



**Docker** is the most popular container engine and runs on Mac, Windows, and Linux.




**Podman** is a new, open-source container engine that runs on Linux.



**Singularity** is a container engine that was designed for HPC systems.

# Container Orchestrators

Orchestrators manage multiple containers (advanced topic)




**Compose** is a simple orchestrator that ships with Docker.




**Kubernetes** is a much more sophisticated orchestrator aimed at deploying complex containerized apps.

# Today's Goals




# Today's Goals

- Learn how to **pull images** from Docker Hub.



# Today's Goals


- Learn how to **pull images** from Docker Hub.
- Learn how to **start a container** interactively.





# Today's Goals

- Learn how to **pull images** from Docker Hub.
- Learn how to **start a container** interactively.
- Learn how to **build your own container images**.



# Pulling Images

# Pulling Images

Docker's official repository of container images:

# Pulling Images

Docker's official repository of container images:

**Docker Hub** (<https://hub.docker.com/>)

# Pulling Images

Docker's official repository of container images:

**Docker Hub** (<https://hub.docker.com/>)

You can find a variety of images for:

# Pulling Images

Docker's official repository of container images:

**Docker Hub** (<https://hub.docker.com/>)

You can find a variety of images for:

- **Linux distributions** (e.g. Ubuntu, Fedora, Manjaro, Clear Linux)

# Pulling Images

Docker's official repository of container images:

**Docker Hub** (<https://hub.docker.com/>)

You can find a variety of images for:

- **Linux distributions** (e.g. Ubuntu, Fedora, Manjaro, Clear Linux)
- **Programming languages** (e.g. Python, Julia, c, C++, Fortran)

# Pulling Images

Docker's official repository of container images:

**Docker Hub** (<https://hub.docker.com/>)

You can find a variety of images for:

- **Linux distributions** (e.g. Ubuntu, Fedora, Manjaro, Clear Linux)
- **Programming languages** (e.g. Python, Julia, c, C++, Fortran)
- **Software** (Jupyter Lab, Tensor Flow, Paraview)





```
$ docker pull fedora:latest
```

```
$ docker pull fedora:latest
```

*pulls the image named fedora from docker hub onto your local computer*

# Starting a Container

# Starting a Container

- To start a container use **docker run**.

# Starting a Container

- To start a container use **docker run**.
- The format for the command is:

```
$ docker run [FLAGS] image:tag command
```

<https://docs.docker.com/engine/reference/run/>



```
$ docker run -ti fedora:latest bash
```



**image**



**\$ docker run -ti fedora:latest bash**

**image**

**run interactively**

**\$ docker run -ti fedora:latest bash**

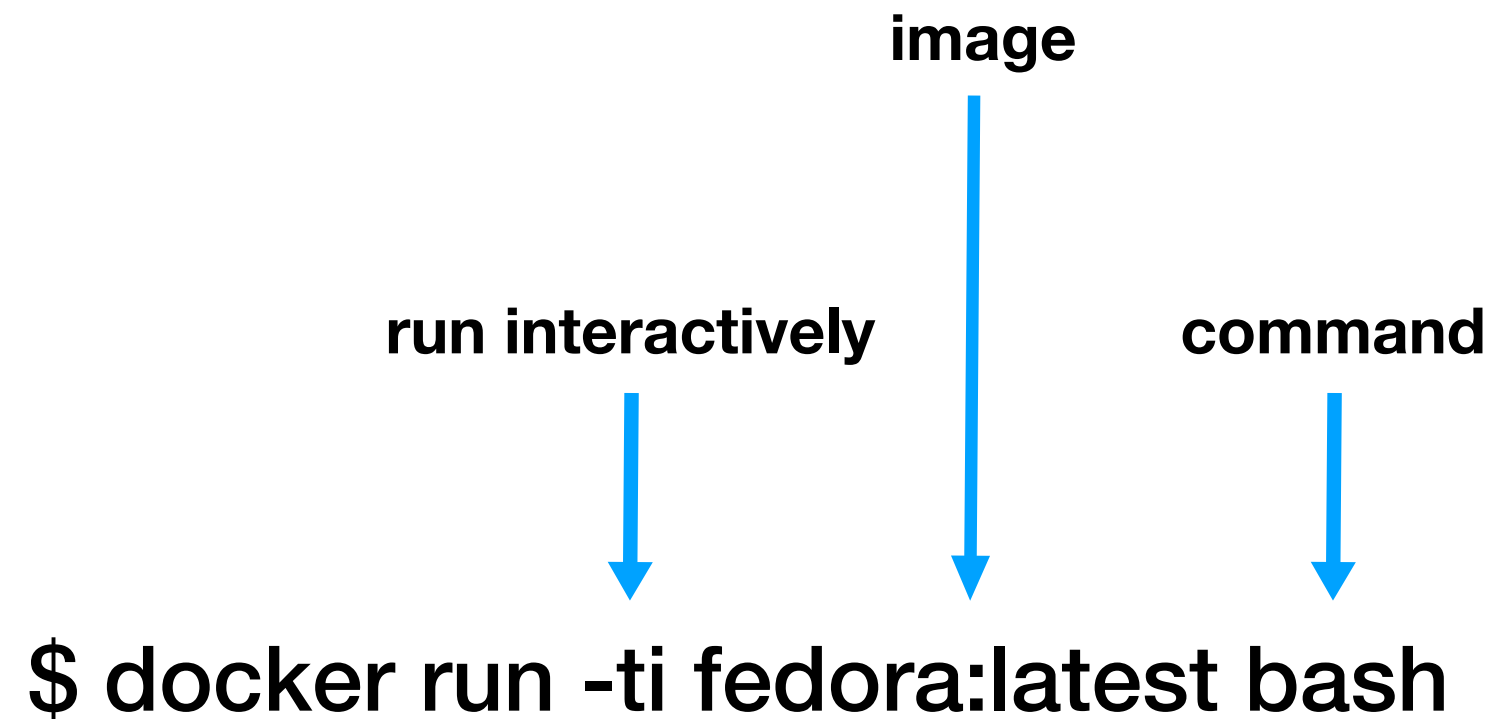


**image**

**run interactively**

**command**

**\$ docker run -ti fedora:latest bash**



*runs the command “bash” interactively using the image fedora:latest*

**image**

**run interactively**

**command**

```
graph TD; image[image] --> fedora; run_interactively[run interactively] --> ti; command[command] --> bash; $ docker run -ti fedora:latest bash
```

**\$ docker run -ti fedora:latest bash**

*runs the command “bash” interactively using the image fedora:latest*

*Let's verify that containers are ephemeral.*

```
$ docker run -ti ubuntu:latest bash
```

*It's easy to try new Linux distributions.*

# Modifying Local Files

# Modifying Local Files

- By default, containers **cannot access** your local files.



# Modifying Local Files

- By default, containers **cannot access** your local files.
- We can grant access by creating a bind mount:

```
$ docker run -v /path/on/host:/path/on/container ...
```

# Modifying Local Files

- By default, containers **cannot access** your local files.

- We can grant access by creating a bind mount:

```
$ docker run -v /path/on/host:/path/on/container ...
```

- This mounts the directory */path/on/host* on the **host** to */path/on/container* in the **container**.



```
$ docker run -v $(pwd):/files -ti fedora:latest bash
```

**binds current directory to /files**



**\$ docker run -v \$(pwd):/files -ti fedora:latest bash**

**binds current directory to /files**



**\$ docker run -v \$(pwd):/files -ti fedora:latest bash**

*runs the command “bash” interactively using the image fedora:latest  
we now also have access to the current directory*

**binds current directory to /files**



**\$ docker run -v \$(pwd):/files -ti fedora:latest bash**

*runs the command “bash” interactively using the image fedora:latest  
we now also have access to the current directory*

*Let's try creating a file. It will also be on the host now.*

# Running Scripts



# Running Scripts

We can run a script from inside the container interactively:

```
$ docker run -v "$(pwd)":/files -ti fedora:latest bash  
[root@c092da88cd6c] $ python3 /files/hello-world.py
```

# Running Scripts

We can run a script from inside the container interactively:

```
$ docker run -v "$(pwd)":/files -ti fedora:latest bash  
[root@c092da88cd6c] $ python3 /files/hello-world.py
```

We can also run it non-interactively:

# Running Scripts

We can run a script from inside the container interactively:

```
$ docker run -v "$(pwd)":/files -ti fedora:latest bash  
[root@c092da88cd6c] $ python3 /files/hello-world.py
```

We can also run it non-interactively:

```
docker run -v "$(pwd)":/files -ti fedora:latest python3 /files/hello-world.py
```

# Running Scripts

We can run a script from inside the container interactively:

```
$ docker run -v "$(pwd)":/files -ti fedora:latest bash  
[root@c092da88cd6c] $ python3 /files/hello-world.py
```

We can also run it non-interactively:

```
docker run -v "$(pwd)":/files -ti fedora:latest python3 /files/hello-world.py
```

**Python does not need to be installed on the host!**

# Building Custom Images

# Building Custom Images

- To build a container image use **docker build**.

# Building Custom Images

- To build a container image use **docker build**.
- You describe the container in a **Dockerfile**.

# Building Custom Images

- To build a container image use **docker build**.
- You describe the container in a **Dockerfile**.
- The format for the command is:

```
$ docker build [FLAGS] -t=image:tag .
```

<https://docs.docker.com/engine/reference/build/>



# Two Build Commands

- `FROM image:tag`  
Allows you to start from an existing container image.
- `RUN command`  
Runs a command during the image build stage.

# A Very Simple Dockerfile

```
FROM fedora:latest
```

```
RUN dnf -y install pip
```

```
# Starts from fedora image
```

```
# install pip
```

**The build context; must contain the Dockerfile.**

**The image name:tag**



```
$ docker build -t myfedora: .
```

*this command will build the image specified by the Dockerfile in the cwd*

*Let's discuss layers while our image builds.*

# Layers

Each line in a Dockerfile creates a new *layer*.

If you don't use the `--no-cache` flag when building, then previously computed layers will be reused.

FROM fedora:latest	# starts from fedora image
RUN dnf -y install pip	# install pip
RUN pip install matplotlib	# install matplotlib

*Let's try running an interactive bash session in our new image.*

# Additional Useful Docker Commands

## Container Management

- Stop a container: `docker stop ID`
- Remove a container: `docker rm ID`
- List running a container: `docker ps`

## Image Management

- Push an image to Dockerhub: `docker push image:tag`
- Remove an Image: `docker rmi image:tag`

## Documentation

<https://docs.docker.com/engine/reference/commandline/docker/>

# Docker Compose

Encode your settings into a yml file

# Docker Compose

## Encode your settings into a yml file

- If you use Docker, the command to start an interactive terminal is:

```
$ docker run -ti -v $(pwd) : /root/files myfedora: pip bash
```

# Docker Compose

## Encode your settings into a yml file

- If you use Docker, the command to start an interactive terminal is:

```
$ docker run -ti -v $(pwd) : /root/files myfedora: pip bash
```

- This is not a simple command.



# Docker Compose

## Encode your settings into a yml file

- If you use Docker, the command to start an interactive terminal is:

```
$ docker run -ti -v $(pwd) : /root/files myfedora: pip bash
```

- This is not a simple command.
- Using Docker Compose, we can encode this into a file.

# Docker Compose

## Encode your settings into a yml file

- If you use Docker, the command to start an interactive terminal is:

```
$ docker run -ti -v $(pwd) : /root/files myfedora:pip bash
```

- This is not a simple command.
- Using Docker Compose, we can encode this into a file.
- We can run the whole command with `docker-compose run`.

# A Simple Compose File

**docker-compose.yml**

```
version: "3.8"
services:
  bash:
    image: myfedora:pip
    volumes:
      - ./root/files
    command: bash
```

**\$ docker-compose run bash**

# If you find containers interesting...

## Here is an even simpler tool



**container job runner** - a simple way to develop and run code in containers.

### Local Development

```
$ cjr shell --here          # start an interactive shell in local directory
$ cjr jupyter:start        # start Jupyter Lab or notebook
```

### Jobs

```
$ cjr job:start --resource=my-server python myscript.py    # run remote job
$ cjr job:start --resource=localhost python myscript.py    # run local job
```

**Thanks for your attention!**